

**Example :** Consider  $p = 5$ . Suppose the set of random integers  $(i_1, i_2, i_3, i_4)$  are  $\{1, 1, 3, 4\}$  respectively. Then, the sequence  $\{(x, i_1, i_2, i_3, i_4, x)\}$  will have matrix representation as

$$\begin{bmatrix} 0 & 1 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

### 3 Encoder-Decoder architecture [10 marks]

You have to train an encoder-decoder model to learn a transformation which takes a sentence and outputs a transformed sentence. Both the sentence and the transformed sentence are 8 characters long and only consist of lowercase English alphabets. You have been provided with “train\_data.csv” which consists of 7000 sentence-transformed sentence examples and “eval\_data.csv” which consists of 2000 sentence-transformed sentence examples. Make note of the following points:

1. You are not supposed to use the sentences provided in “eval\_data.csv” for training the model. You can use it to judge overfitting or for early stopping.
2. You are free to use whatever architecture you wish, i.e., LSTM vs RNN vs GRU, num of hidden units, num of layers, etc (**Attention networks and Transformers are not allowed**).
3. You are free to train your model however you wish (Early stopping, dropout, optimizers, regularization, number of epochs, etc).
4. You are free to explore any other approaches to improve your model’s performance as well (better embeddings, data preprocessing etc).
5. **You may want to explore the “teacher forcing” technique for training RNNs and embedding layers in pytorch/tensorflow.**
6. You do not need to provide any kind of analysis for this problem, but you are encouraged to conduct analysis of your own to find ways to improve your model.
7. **If something is present in your code, you should know how it works. You should have some idea of what the particular function does and how it achieves it. Saying you just copied it from the internet will not suffice.** During evals if the TAs determine you have simply copied some part of the code from the internet without understanding it, you may be penalized for this question.

### 3.1 Evaluation

The testing for this question will be automated, the submission instructions along with a checker script will be released on Oct 2<sup>nd</sup>. The final evaluation will be run on GPU/CUDA, so make sure you write your code accordingly. The max training time for the model will be 2 hours, after which it will be terminated (early stopping might be useful).

Note: The final submissions for Q3 will be a python script, not a jupyter notebook. You can use it for coding (such as google colab), but you need to convert it to a python script (google how to do so) and submit it for final evaluation.

### 3.2 Marks distribution

- 4 marks will be based on the performance of your model on “eval\_data”.
- 6 marks will be based on the performance of your model on unseen data consisting of 1000 examples.

A prediction will be considered correct if 6 or more characters are identified correctly. A prediction will be considered partially correct if 5 characters are predicted correctly. You will be awarded 1 point for every correct prediction and 0.5 points for every partially correct prediction.

Your marks for “eval\_data” will be calculated by the formula:

$$\text{eval\_marks} = \min(4, \frac{\text{score}}{1600} \times 4)$$

Your marks for the unseen data will be calculated by the formula:

$$\text{marks} = \min(6, \frac{\text{score}}{800} \times 6)$$

Marks will be rounded to the nearest .5 marks.